

# Final Project

Predict credit score  
using multi-algorithm bagging

Statistical Learning|張致語|108021186

# Overview

- Inspiration
- Data and Algorithm
- Sub-model1: SVM classifier
- Sub-model2: Tree classifier
- Sub-model3: Logistic regression
- Sub-model4: QDA
- Aggregate Process and comparison
- Conclusion

# Inspiration

- Paper 'Improving Bagging Performance through Multi-Algorithm Ensembles' connects the diversity with correlation and further shows that with higher diversity will improve accuracy in classification problems.



# Data and algorithm

- Source of data



DEEPAK SAI PENDYALA · UPDATED 19 DAYS AGO



New Notebook

Download (4 MB)



## credit\_score\_classification\_processed

Given a person's credit-related information, build a machine learning model that



- Link:  
<https://www.kaggle.com/datasets/deepaksaipendyala/credit-score-classification-processed>

# Data and algorithm

- Source of data

Data update: 2022/12/18

Dimension of data: 25 variables with 79805 rows

- Variables

X <int>	Month <int>	Age <dbl>	Occupation <int>	Annual_Income <dbl>	Monthly_Inhand_Salary <dbl>	Num_Bank_Accounts <dbl>
1	2	23.00000	13	19114.12	4194.171	3
2	6	34.42982	13	19114.12	4194.171	3
3	0	23.00000	13	19114.12	4194.171	3
4	7	23.00000	13	19114.12	1824.843	3
5	5	23.00000	13	19114.12	4194.171	3
6	4	23.00000	13	19114.12	1824.843	3

Categorical variable

# Data and algorithm

- Variables

Num_Credit_Card <dbl>	Interest_Rate <dbl>	Num_of_Loan <dbl>	Type_of_Loan <int>	Delay_from_due_date <int>
4	3	4	128	-1
4	3	4	128	3
4	3	4	128	5
4	3	4	128	6
4	3	4	128	8
4	3	4	128	3

Categorical variable

Categorical variable

# Data and algorithm

- Variables

Num_of_Delayed_Payment <dbl>	Changed_Credit_Limit <int>	Num_Credit_Inquiries <dbl>	Credit_Mix <int>	Outstanding_Debt <dbl>
30.92334	1112	4	1	809.98
7.00000	4238	4	1	809.98
4.00000	3764	4	1	809.98
30.92334	1112	4	1	809.98
4.00000	4163	4	1	809.98
8.00000	1112	4	1	809.98

Categorical variable

# Data and algorithm

## Variables

Categorical variable

Credit_Utilization_Ratio <dbl>	Credit_History_Age <int>	Payment_of_Min_Amount <int>	Total_EMI_per_month <dbl>
31.94496	0	1	49.57495
28.60935	267	1	49.57495
31.37786	268	1	49.57495
24.79735	269	1	49.57495
27.26226	270	1	49.57495
22.53759	271	1	49.57495

Total_EMI_per_month <dbl>	Amount_invested_monthly <dbl>	Payment_Behaviour <int>	Monthly_Balance <dbl>	Credit_Score <int>
49.57495	118.28022	3	284.6292	2
49.57495	81.69952	4	331.2099	2
49.57495	199.45807	5	223.4513	2
49.57495	41.42015	1	341.4892	2
49.57495	62.43017	6	340.4792	2
49.57495	178.34407	5	244.5653	2

Categorical variable

Dependent variable



# Data and algorithm

- Pre-processing

```
age_out=which(data$Age%%1≠0)
card_out=which(data$Num_Credit_Card%%1≠0)
account_out=which(data$Num_Bank_Accounts%%1≠0)
inquire_out=which(data$Num_Credit_Inquiries%%1≠0)
delay_out=which(data$Num_of_Delayed_Payment%%1≠0 | data$Num_of_Delayed_Payment<0)
date_out=which(data$Delay_from_due_date<0)
data_out=Reduce(union,list(age_out,card_out,account_out,inquire_out,delay_out,date_out))
data=data[-data_out,]
```

Remove NA's and wrong collected data

```
data$Occupation=factor(data$Occupation)
data$Type_of_Loan=factor(data$Type_of_Loan)
data$Interest_Rate=factor(data$Interest_Rate)
data$Credit_Score=factor(data$Credit_Score)
data$Payment_of_Min_Amount=factor(data$Payment_of_Min_Amount)
data$Payment_Behaviour=factor(data$Payment_Behaviour)
```

Change categorical variables into factor

## Data and algorithm

- EDA: summary of data and special findings

1. The response is very unbalanced

Credit\_Score

0:17902

1:33964

2:11554

We need to balanced the training data to avoid the overfitting.

2. The range of month balance and income are very large

Annual_Income		Monthly_Balance	
Min.	: 7006	Min.	: 0.0886
1st Qu.:	19645	1st Qu.:	272.0805
Median :	38107	Median :	341.6435
Mean :	179116	Mean :	406.2164
3rd Qu.:	73850	3rd Qu.:	474.1676
Max.	:24198062	Max.	:1602.0405

## Data and algorithm

- EDA: summary of data and special findings




3. There are some categorical variables that exist more than 32 categories

We need to use modules other than tree when fitting classification tree.

Occupation		Type_of_Loan	
12	: 4456	6260	: 7397
7	: 4192	3463	: 919
4	: 4036	4878	: 859
13	: 4024	5591	: 851
1	: 4013	684	: 845
0	: 4000	1410	: 838
(Other)	: 38682	(Other)	: 51694

# Data and algorithm

4. There are a lot difference in income, credit card and loan of different credit score



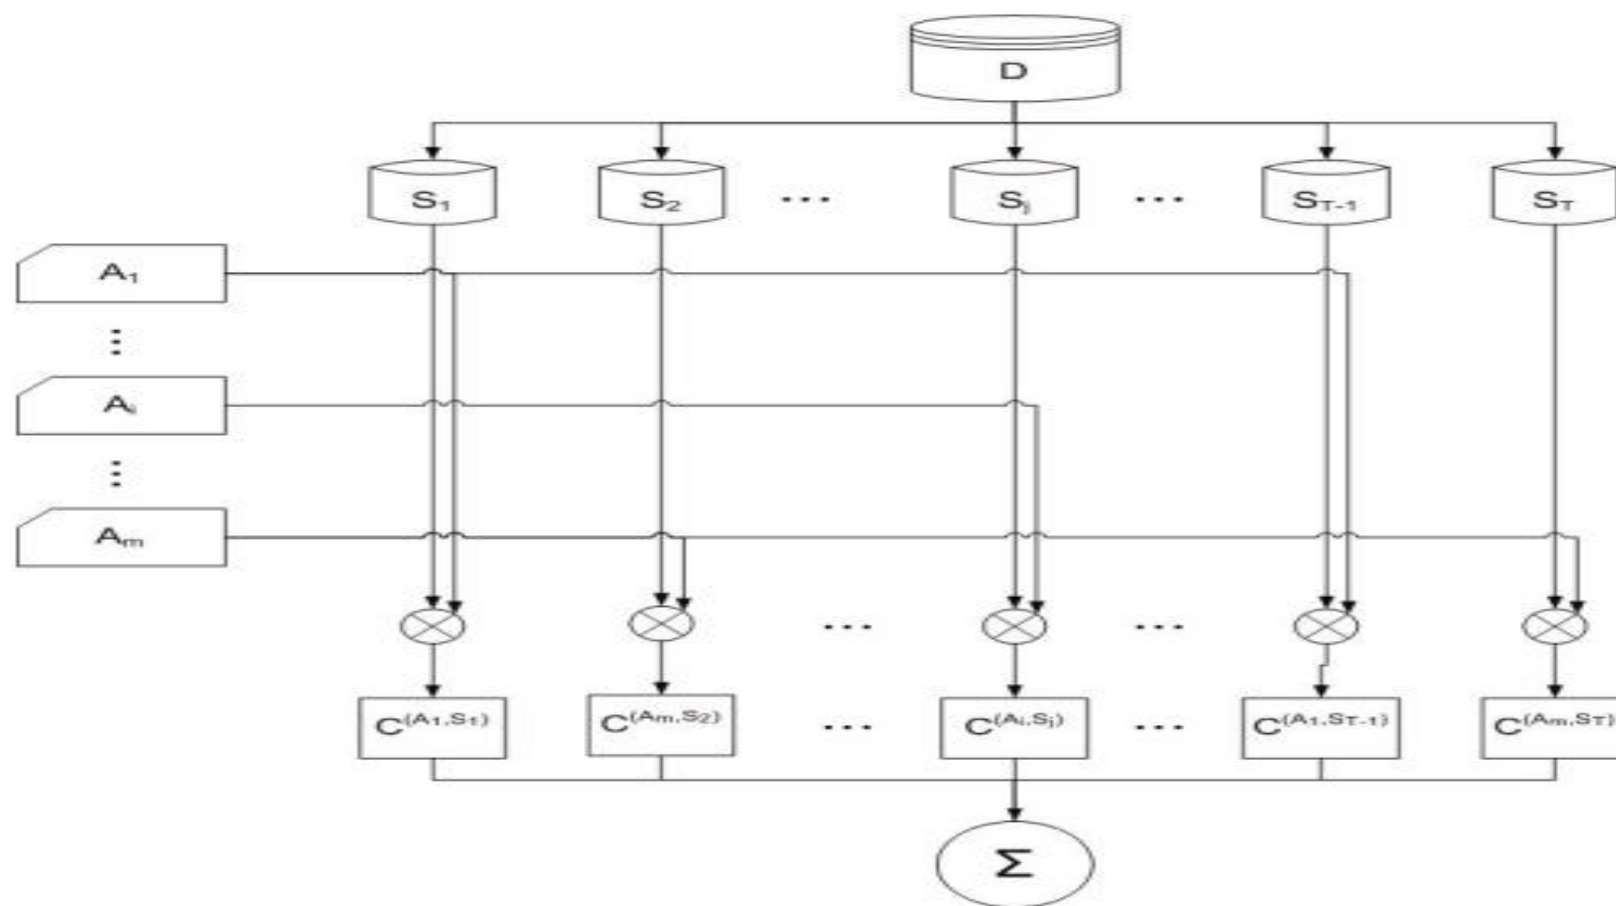
# Data and algorithm

- Algorithm

1. Train data: contains 2400 data with each label are equality counted
2. Test data: 61003 rows of data
3. Variables: 22 independent variables and one dependent variables
4. Algorithm bag:

SVM model, Classification tree, Logistic regression, LDA

# Data and algorithm



# Data and algorithm

- Algorithm

**Algorithm 1** (Framework for bagging with multi-algorithm ensembles)

1. Initialize  $C^* = \emptyset$
2. For  $j = 1$  to  $T$
3.    $S_j \leftarrow \text{Bootstrap}(D, p = 1)$
4.   Select an algorithm  $A_i$  from the given bag of algorithms  $M$
5.    $C_i \leftarrow \mathbf{C}(S_i, A_i)$ , where  $\mathbf{C}$  is a procedure, e.g. API, to create a classifier
6.    $C^* \leftarrow C^* \cup C_i$
7. End for

# Sub-model1: SVM classifier

Choose 400 data as validation set to choose best cost and gamma

Fitting the model

Perform 10 classifiers to see the divergence

```
library(e1071)
set.seed(108021186)
tune.out <- tune(svm, Credit_Score ~ ., data = train[1:400,], kernel = "linear",
  ranges = list(cost=c(0.1,1,10,100,1000,10000,100000)))
summary(tune.out)
```
```

R Console

data.frame  
1 x 1

data.frame  
7 x 3

| cost<br><dbl> | error<br><dbl> | dispersion<br><dbl> |
|---------------|----------------|---------------------|
| 1e-01         | 0.3600         | 0.06687468          |
| 1e+00         | 0.4350         | 0.06032320          |
| 1e+01         | 0.4850         | 0.07564537          |
| 1e+02         | 0.4975         | 0.06503204          |
| 1e+03         | 0.4950         | 0.06324555          |
| 1e+04         | 0.4950         | 0.06324555          |
| 1e+05         | 0.4950         | 0.06324555          |



# Sub-model1: SVM classifier

Choose 400 data as validation set to choose best cost and gamma



Fitting the model



Perform 10 classifiers to see the divergence

```
library(e1071)
set.seed(108021186)
svm_result=1:61003
for (i in 1:20){
  boot=train[sample(1:2400,2400,replace=T),]
  svmfit<-svm(Credit_Score~.,data=boot,kernel="linear",
  cost=0.1,scale = FALSE)
  svm_result=cbind(svm_result,predict(svmfit,test[, -23]))
  print(i)
}
svm_result_mod=svm_result[, -1]-1
```

# Sub-model1: SVM classifier

Choose 400 data as validation set to choose best cost and gamma



Fitting the model



Perform 10 classifiers to see the divergence

Aggregate 20 classifiers using simple mean

Accruacy:48.27%

```
result=function(mat){  
  resul=NULL  
  for(i in 1:nrow(mat)){  
    a0=sum(mat[i,]==0)  
    a1=sum(mat[i,]==1)  
    a2=sum(mat[i,]==2)  
    res=which.max(c(a0,a1,a2))  
    resul[i]=res-1  
  }  
  return(resul)  
}  
``
```

```
``{r}  
res=result(svm_result_mod)  
table(res,test$Credit_Score)  
``
```

| res | 0    | 1     | 2    |
|-----|------|-------|------|
| 0   | 5578 | 5381  | 531  |
| 1   | 9837 | 19397 | 5744 |
| 2   | 1680 | 8382  | 4473 |

# Sub-model1: SVM classifier

Choose 400 data as  
validation set to choose  
best cost and gamma

Fitting the model

Perform 10 classifiers to  
see the divergence

Correlation matrix with first 10 classifiers


|        |        |        |        |        |        |        |        |        |        |
|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|
| 1.000  | 0.506  | 0.199  | 0.460  | -0.235 | 0.220  | -0.096 | -0.378 | 0.193  | 0.288  |
| 0.506  | 1.000  | 0.278  | 0.614  | -0.560 | 0.446  | -0.337 | -0.648 | 0.158  | 0.374  |
| 0.199  | 0.278  | 1.000  | 0.062  | -0.418 | -0.027 | 0.236  | -0.488 | 0.539  | -0.091 |
| 0.460  | 0.614  | 0.062  | 1.000  | -0.316 | 0.490  | -0.421 | -0.319 | -0.154 | 0.471  |
| -0.235 | -0.560 | -0.418 | -0.316 | 1.000  | -0.297 | 0.003  | 0.703  | -0.213 | -0.162 |
| 0.220  | 0.446  | -0.027 | 0.490  | -0.297 | 1.000  | -0.452 | -0.215 | -0.193 | 0.372  |
| -0.096 | -0.337 | 0.236  | -0.421 | 0.003  | -0.452 | 1.000  | -0.290 | 0.501  | -0.430 |
| -0.378 | -0.648 | -0.488 | -0.319 | 0.703  | -0.215 | -0.290 | 1.000  | -0.453 | -0.019 |
| 0.193  | 0.158  | 0.539  | -0.154 | -0.213 | -0.193 | 0.501  | -0.453 | 1.000  | -0.170 |
| 0.288  | 0.374  | -0.091 | 0.471  | -0.162 | 0.372  | -0.430 | -0.019 | -0.170 | 1.000  |

## Sub-model2: Classification tree

Choose 400 data as  
validation set to choose  
best cost and gamma

Fitting the model

Perform 10 classifiers to  
see the divergence



## Sub-model2: Classification tree

Choose 400 data as validation set to choose best cost and gamma

Fitting the model

Perform 10 classifiers to see the divergence

Choose max.depth=7

```
library(rpart)
set.seed(108021186)
tree_result=1:61003
for (i in 1:20){
  boot=train1[sample(1:2400,2400,replace=T),]
  treefit<-rpart(Credit_Score~.,data = boot, maxdepth=10, cp=0)
  tree_result=cbind(tree_result,predict(treefit,test1[, -1],type='class'))
  print(i)
}
tree_result_mod=tree_result[, -1]-1
```

Aggregate 20 classifiers using simple mean

Accruacy:67.40%

| agg_tree | 0     | 1     | 2    |
|----------|-------|-------|------|
| 0        | 12139 | 6429  | 508  |
| 1        | 2853  | 20348 | 1610 |
| 2        | 2103  | 6383  | 8630 |

## Sub-model2: Classification tree

Choose 400 data as  
validation set to choose  
best cost and gamma



Fitting the model



Perform 10 classifiers to  
see the divergence

Correlation matrix with first 10 classifiers

|       |       |       |       |       |       |       |       |       |       |
|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| 1.000 | 0.356 | 0.342 | 0.330 | 0.310 | 0.321 | 0.330 | 0.329 | 0.334 | 0.357 |
| 0.356 | 1.000 | 0.404 | 0.377 | 0.388 | 0.324 | 0.404 | 0.409 | 0.407 | 0.409 |
| 0.342 | 0.404 | 1.000 | 0.369 | 0.387 | 0.353 | 0.394 | 0.375 | 0.396 | 0.395 |
| 0.330 | 0.377 | 0.369 | 1.000 | 0.339 | 0.339 | 0.388 | 0.368 | 0.361 | 0.347 |
| 0.310 | 0.388 | 0.387 | 0.339 | 1.000 | 0.359 | 0.379 | 0.382 | 0.364 | 0.366 |
| 0.321 | 0.324 | 0.353 | 0.339 | 0.359 | 1.000 | 0.362 | 0.319 | 0.365 | 0.328 |
| 0.330 | 0.404 | 0.394 | 0.388 | 0.379 | 0.362 | 1.000 | 0.373 | 0.395 | 0.399 |
| 0.329 | 0.409 | 0.375 | 0.368 | 0.382 | 0.319 | 0.373 | 1.000 | 0.381 | 0.393 |
| 0.334 | 0.407 | 0.396 | 0.361 | 0.364 | 0.365 | 0.395 | 0.381 | 1.000 | 0.388 |
| 0.357 | 0.409 | 0.395 | 0.347 | 0.366 | 0.328 | 0.399 | 0.393 | 0.388 | 1.000 |

## Sub-model3: Logistic regression

Fitting the model

Perform 10 classifiers to  
see the divergence

```
library(nnet)
set.seed(108021186)
lr_result=1:61003
for (i in 1:20){
  boot=train[sample(1:2400,2400,replace=T),]
  lrfit<-nnet::multinom(Credit_Score ~., data = train)
  lr_result=cbind(lr_result,predict(lrfit,test[, -20],type='class'))
  print(i)
}
lr_result_mod=lr_result[, -1]-1
```

Accruacy:60.16%

| res | 0    | 1     | 2    |
|-----|------|-------|------|
| 0   | 9363 | 6213  | 310  |
| 1   | 4840 | 18498 | 1597 |
| 2   | 2892 | 8449  | 8841 |

## Sub-model3: Logistic regression

Fitting the model



Perform 10 classifiers to  
see the divergence

|       | [,1]  | [,2]  | [,3]  | [,4]  | [,5]  | [,6]  | [,7]  | [,8]  | [,9]  | [,10] |
|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| [1,]  | 1.000 | 0.810 | 0.802 | 0.830 | 0.839 | 0.801 | 0.839 | 0.829 | 0.828 | 0.828 |
| [2,]  | 0.810 | 1.000 | 0.793 | 0.824 | 0.825 | 0.800 | 0.814 | 0.791 | 0.833 | 0.827 |
| [3,]  | 0.802 | 0.793 | 1.000 | 0.824 | 0.805 | 0.787 | 0.814 | 0.777 | 0.803 | 0.801 |
| [4,]  | 0.830 | 0.824 | 0.824 | 1.000 | 0.833 | 0.840 | 0.834 | 0.840 | 0.868 | 0.833 |
| [5,]  | 0.839 | 0.825 | 0.805 | 0.833 | 1.000 | 0.820 | 0.839 | 0.814 | 0.841 | 0.867 |
| [6,]  | 0.801 | 0.800 | 0.787 | 0.840 | 0.820 | 1.000 | 0.839 | 0.792 | 0.845 | 0.837 |
| [7,]  | 0.839 | 0.814 | 0.814 | 0.834 | 0.839 | 0.839 | 1.000 | 0.795 | 0.855 | 0.842 |
| [8,]  | 0.829 | 0.791 | 0.777 | 0.840 | 0.814 | 0.792 | 0.795 | 1.000 | 0.832 | 0.802 |
| [9,]  | 0.828 | 0.833 | 0.803 | 0.868 | 0.841 | 0.845 | 0.855 | 0.832 | 1.000 | 0.849 |
| [10,] | 0.828 | 0.827 | 0.801 | 0.833 | 0.867 | 0.837 | 0.842 | 0.802 | 0.849 | 1.000 |



## Sub-model4: QDA classifier

Fitting the model

Perform 10 classifiers to see the divergence

```
library(MASS)
set.seed(108021186)
qda_result=1:61003
for (i in 1:20){
  boot=train[sample(1:2400,2400,replace=T),]
  qdafit<-qda(Credit_Score[,boot])
  qda_result=cbind(qda_result,predict(qdafit,test[,20])$class)
  print(i)
}
qda_result_mod=qda_result[,-1]-1
```

Accuracy:51.56%

| agg_qda | 0     | 1     | 2    |
|---------|-------|-------|------|
| 0       | 12749 | 14820 | 855  |
| 1       | 1454  | 10043 | 1228 |
| 2       | 2892  | 8297  | 8665 |

## Sub-model4: QDA classifier

Fitting the model



Perform 10 classifiers to  
see the divergence

|       |       |       |       |       |       |       |       |       |       |
|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|
| 1.000 | 0.679 | 0.641 | 0.658 | 0.672 | 0.631 | 0.677 | 0.670 | 0.678 | 0.661 |
| 0.679 | 1.000 | 0.708 | 0.742 | 0.650 | 0.700 | 0.675 | 0.693 | 0.639 | 0.684 |
| 0.641 | 0.708 | 1.000 | 0.709 | 0.717 | 0.691 | 0.717 | 0.694 | 0.712 | 0.737 |
| 0.658 | 0.742 | 0.709 | 1.000 | 0.648 | 0.668 | 0.685 | 0.691 | 0.655 | 0.690 |
| 0.672 | 0.650 | 0.717 | 0.648 | 1.000 | 0.671 | 0.736 | 0.756 | 0.774 | 0.717 |
| 0.631 | 0.700 | 0.691 | 0.668 | 0.671 | 1.000 | 0.698 | 0.703 | 0.642 | 0.676 |
| 0.677 | 0.675 | 0.717 | 0.685 | 0.736 | 0.698 | 1.000 | 0.769 | 0.755 | 0.708 |
| 0.670 | 0.693 | 0.694 | 0.691 | 0.756 | 0.703 | 0.769 | 1.000 | 0.773 | 0.710 |
| 0.678 | 0.639 | 0.712 | 0.655 | 0.774 | 0.642 | 0.755 | 0.773 | 1.000 | 0.737 |
| 0.661 | 0.684 | 0.737 | 0.690 | 0.717 | 0.676 | 0.708 | 0.710 | 0.737 | 1.000 |

# Aggregation process and comparison

- Comparison of different sub-model

| model               | Single accuracy | Aggregated accuracy | System time |
|---------------------|-----------------|---------------------|-------------|
| SVM model           | 54.00%          | 48.27%              | 123.09      |
| Tree model          | 61.67%          | 67.40%              | 0.16        |
| Logistic regression | 60.16%          | 60.16%              | 0.41        |
| QDA model           | 54.09%          | 51.56%              | 0.06        |

# Aggregation process and comparison

- Aggregation process

```
set.seed(108021186)
setting=sample(1:80,20)
mult_result=cbind(svm_result_mod,tree_result_mod,lr_result_mod,qda_result_mod)[,setting]
agg_mult=result(mult_result)
table(agg_mult,test$Credit_Score)
mean(agg_mult==test$Credit_Score)
round(cor_dist(mult_result[,1:10],test$Credit_Score),3)
```
```

```
agg_mult      0      1      2
      0 12267  8220  400
      1  2039 17302 1502
      2  2789  7638 8846
[1] 0.6297231
```

# Aggregation process and comparison

- Aggregation process

Correlation matrix

1.000	0.055	0.332	0.286	0.021	0.041	0.075	0.053	0.316	0.288
0.055	1.000	0.317	0.168	0.347	-0.090	0.365	0.339	0.319	0.309
0.332	0.317	1.000	0.237	0.320	-0.126	0.360	0.314	0.756	0.686
0.286	0.168	0.237	1.000	0.191	0.003	0.185	0.143	0.277	0.216
0.021	0.347	0.320	0.191	1.000	-0.041	0.398	0.345	0.310	0.313
0.041	-0.090	-0.126	0.003	-0.041	1.000	-0.073	-0.047	-0.117	-0.177
0.075	0.365	0.360	0.185	0.398	-0.073	1.000	0.361	0.353	0.362
0.053	0.339	0.314	0.143	0.345	-0.047	0.361	1.000	0.309	0.317
0.316	0.319	0.756	0.277	0.310	-0.117	0.353	0.309	1.000	0.687
0.288	0.309	0.686	0.216	0.313	-0.177	0.362	0.317	0.687	1.000

# Conclusion

- Discussion

1. Lower correlation between classifiers would be more reasonable to aggregate them using simple mean.
2. With multi-algorithm, it seems to lower the performance with just using tree model. But it become more legitimate.
3. We also need to pay attention to time cost while using multi-algorithm bagging model.